

Approximate Bayesian Computation

Dom Hutchinson

December 29, 2020

Contents

1	Intro to ABC	1
1.1	Decisions	2
2	ABC Algorithms	3
3	Semi-Automatic ABC	5

1 Intro to ABC

Definition 1.1 - Approximate Bayesian Computation (ABC)

Approximate Bayesian Computation (ABC) is a family of computational methods for estimating the posterior of model parameters for *Generative Models*. *Generative Models* are models which can be simulated from but we do not have an explicit definition for their posterior $f_{\mathcal{G}}(x|\theta)$ (eg most IRL systems).

Proposition 1.1 - Motivating Idea[1]

Consider a set of observations^{[1][2]} $\mathbf{y} := (y_1, \dots, y_n)$ where each $y_i \in \mathbb{R}^m$ is high dimensional. Let $s(\cdot) : \mathbb{R}^m \rightarrow \mathbb{R}^p$ be a mapping (known as a *Summary Statistic*) from the observed data to some lower dimension p .

ABC aims to infer the joint distribution of parameters θ and general summary statistics \mathbf{s} , given the observed summary statistics $\mathbf{s}_{obs} := s(\mathbf{y})$

$$p_{\epsilon}(\theta, \mathbf{s}|\mathbf{s}_{obs}) \propto \underbrace{\pi_0(\theta)}_{\text{Prior}} \underbrace{f(s|\theta)}_{\text{Likelihood}} K_{\epsilon}(\|\mathbf{s} - \mathbf{s}_{obs}\|)$$

where $K_{\epsilon}(\cdot)$ is a kernel function with scaling parameter ϵ and $\|\cdot\|$ is a distance measure (e.g. Euclidean).^[3]

From this joint distribution the posterior for parameters θ , given the observed summary statistics $\mathbf{s}_{obs} := s(\mathbf{y})$, can be calculated as

$$p_{\epsilon}(\theta|\mathbf{s}_{obs}) = \int p_{\epsilon}(\theta, \mathbf{s}|\mathbf{s}_{obs}) d\mathbf{s}$$

^[1]From a *Generative Model*

^[2]Generally these observations are ordered in some way (by the variables of the system) so can be considered a sequence $\{y_t\}_{t \in T}$ where T specifies the variable values in each epoch.

^[3]The likelihood $f(s|\theta)$ is the only one of these features which is not specified by the user, and thus what we need to “learn” it.

Monte-Carlo Algorithms can be used to sample from this posterior $p_\epsilon(\boldsymbol{\theta}|\mathbf{s}_{obs})$ without having to explicitly state the likelihood $f(\mathbf{s}|\boldsymbol{\theta})$. *Numerical Integration* methods can then be used to evaluate the integral^[4].

Proposition 1.2 - Setup of ABC

To perform ABC we typically have/define the following features

- A set of observations from a *Generative Model* $\mathbf{y} := (y_1, \dots, y_n)$ where each y_i is high-dimensional.
- A map $s(\cdot)$ which maps the high-dimensional observed data to a lower dimension.
- A theorised model with posterior pdf $f_{\mathcal{T}}(\cdot|\boldsymbol{\theta}, \mathbf{x})$ where $\boldsymbol{\theta}$ are the parameters which we wish to fit to the *Generative Model* using ABC.
- A prior $\pi_0(\cdot)$ for the parameters $\boldsymbol{\theta}$.
- A kernel $K_\epsilon(\cdot)$ and a distance measure $\|\cdot\|$.

1.1 Decisions

Remark 1.1 - Decisions

When implementing ABC there are several decisions to make, including:

- What theorised model $f(\cdot|\boldsymbol{\theta}, \mathbf{x})$ to use.
- What kernel $K_\epsilon(\cdot)$ to use.
- What summary statistics $s(\cdot)$ to use.
- Do we even need summary statistics?
- How long to sample for?

Proposition 1.3 - Kernels $K_\epsilon(\cdot)$

A *Kernel* is used to determine with what probability to accept a sample, given it is a certain distance away from observed data. Here are some common kernels

- *Uniform Kernel* $K_\epsilon(\|\mathbf{s} - \mathbf{s}_{obs}\|) := \mathbb{1}\{\|\mathbf{s} - \mathbf{s}_{obs}\| \leq \epsilon\}$ which accepts simulated values if they are within ϵ of observed data.
- *Epanechnikov Kernel* $K_\epsilon(\|\mathbf{s} - \mathbf{s}_{obs}\|) := \frac{3}{4\epsilon} \left(1 - \left(\frac{\|\mathbf{s} - \mathbf{s}_{obs}\|}{\epsilon}\right)^2\right)$ for $\|\mathbf{s} - \mathbf{s}_{obs}\| \leq \epsilon$
- *Gaussian Kernel* $K(\|\mathbf{s} - \mathbf{s}_{obs}\|) := \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}\|\mathbf{s} - \mathbf{s}_{obs}\|^2}$

Proposition 1.4 - Summary Statistics $s(\cdot)$

See `SummaryStatisticSelection.pdf`

Proposition 1.5 - How long to sample for

The algorithm given in **Proposition 1.3** runs the algorithm until a sufficiently large sample has been produced. This is not ideal as the algorithm will run for an unknown period of time and is dependent upon the kernel $K_\epsilon(\cdot)$ which has been defined.

Alternatively, all simulated values could be kept and then all but the best $M^{[5]}$ are discarded.

^[4]See *Monte-Carlo Integration methods: Uniform Sampling, Importance Sampling*

^[5] M closest to \mathbf{s}_{obs} .

2 ABC Algorithms

Proposition 2.1 - ABC Algorithm - Rejection

Consider the setup in Proposition 1.2. Here is a simple, online algorithm for ABC

- i). Sample a set of parameters from the prior $\tilde{\theta}_t \sim \pi_0(\theta)$.
- ii). Sample from the theorised model using these sampled parameters

$$\mathbf{y}_t \sim f_{\mathcal{T}}(\mathbf{y}|\tilde{\theta}_t)$$

- iii). Calculate the summary statistic values for the sampled values $\mathbf{s}_t = \mathbf{s}(\mathbf{y}_t)$.
- iv). Reject the sample summary statistic value \mathbf{s}_t with probability $K_{\epsilon}(\|\mathbf{s}_t - \mathbf{s}_{obs}\|)$ where $\mathbf{s}_{obs} = \mathbf{s}(\mathbf{y})$.
- v). Repeat steps i)-iii) until a total of M simulated values have been accepted.

Our final sample contains a set of M summary statistics, along with the parameter values θ and variable-space points \mathbf{x} , which produced them. This data can be used to approximate the posterior for the parameter values.

Remark 2.1 - ABC-SMC

The idea behind *ABC-Sequential Monte Carlo* is to initially use a kernel with a large acceptance range to produce a sample of parameters which v. roughly approximate of the posterior. And then to refine this sample to improve the posterior, by resampling and tightening the kernel.

Proposition 2.2 - ABC Algorithm - SMC[2][3][4]

Consider the setup in Proposition 1.2. Here is a *Sequential Monte-Carlo* algorithm for ABC

- *Initialisation* - Choose a set of *scaling parameters* $\{\epsilon_1, \dots, \epsilon_T\}$ which are increasingly tight^[6]

$$\{\epsilon_1, \dots, \epsilon_T\} \text{ where } \epsilon_1 > \dots > \epsilon_T$$

Set N to be the number of sets of parameters to sample, choose some summary statistics $s(\cdot)$ and observe y_{obs} from the true model.

- *Initial Sampling Step* - $t = 1$.^{[7][8]}

- i). Sample a set of parameters $\tilde{\theta}_{1,i}$

$$\tilde{\theta}_1 \sim \pi_0(\theta)$$

- ii). Observe the theorised model \mathcal{T} with these sampled parameters

$$\mathbf{y}_{1,i} \sim f_{\mathcal{T}}(\mathbf{y}|\tilde{\theta}_{1,i})$$

- iii). If $K_{\epsilon_1}(\|s(\mathbf{y}_{1,i}) - s(\mathbf{y}_{obs})\|)$:^[9]

- * Store the sampled parameters $\tilde{\theta}_{1,i}$ in set Θ_1 .
- * Set $w_{1,i} = \frac{1}{N}$ and increment i .

^[6]There are automatic methods where ϵ_{t+1} is calculated at the end of step t st convergence is encouraged.

^[7]Here a set of possible parameters Θ_1 is generated.

^[8]It is possible for the same set of parameters to appear in Θ_t multiple times.

^[9]Sample parameters are accepted by the kernel.

* If $i == N$: move to *Resampling Step*.

- *Resampling Step* - $t = 2, \dots, T$.^[10]

i). Sample $\tilde{\theta}_{t,i}$ from set Θ_{t-1} with probability $w_{t-1,j}$. The weights form a KDE.

$$\mathbb{P}(\theta_i^* = \tilde{\theta}_j) = w_{t-1,j} \text{ for } \tilde{\theta}_j \in \Theta_{t-1}$$

ii). Perturb the sample parameters slightly using a *Perturbance Kernel* K^* to get a slightly different parameter set $\theta_{t,i}^*$

$$\theta_{t,i}^* \sim K^*(\theta|\tilde{\theta}_{t,i})$$

iii). If $\theta_{t,i}^*$ is impossible under the prior:^[11] return to i).

iv). Observe the theorised model \mathcal{T} with these perturbed parameters

$$\mathbf{y}_{t,i} \sim f_{\mathcal{T}}(\mathbf{y}|\theta_{t,i}^*)$$

v). If not $K_{\epsilon_t}(\|s(\mathbf{y}_{t,i}) - s(\mathbf{y}_{obs})\|)$:^[12] return to i).

vi). Add $\theta_{t,i}^*$ to Θ_t and assign it weight $w_{t,i}$

$$\tilde{w}_{t,i} = \frac{\pi_0(\theta_{t,i}^*)}{\sum_{j=1}^N w_{t-1,j} \mathbb{P}(K_t(\theta|\tilde{\theta}_{t,i}) = \theta_{t,i}^*)}$$
^[13]

vii). If $|\Theta_t| < N$: increment i and return to i).

viii). Normalise weights

$$w_{t,i} = \frac{\tilde{w}_{t,i}}{\sum_{i=1}^N \tilde{w}_{t,i}}$$

ix). Increment t .

Code for this can be found at <https://stats.stackexchange.com/a/328384>.

Proposition 2.3 - ABC Algorithm - MCMC^[5]

Consider the setup in Proposition 1.2. Here is a *Markov Chain Monte-Carlo* algorithm for ABC

- *Initialisation* - Observe y_{obs} from the true model, choose some summary statistics $s(\cdot)$, define the length of the chain N ^[14], a theorised model T with pdf $f_{\mathcal{T}}(\cdot|\theta)$, define a prior for the priors $\pi_0(\theta)$ and a pertubance kernel $K^*(\cdot)$ ^[15].
- *Starting Sample*

i). Sample $\tilde{\theta}_0 \sim \pi(\theta)$.

^[10]Here, we determine weightings for the sets of parameters found in the previous step. With the weightings representing posterior probabilities.

^[11]ie $\pi_0(\theta_{t,i}^*) = 0$

^[12]ie the observation is reject by the kernel

^[13]Weighted sum of the probability this set of parameters $\theta_{t,i}^*$ could have been observed under another one of the previous samples.

^[14]This is used for the termination condition, but other termination conditions are common. Most are based around spotting convergence.

^[15]e.g. Add some gaussian noise.

- ii). Observe y from the theorised model \mathcal{T} using these sampled parameter θ_0 .

$$y \sim f_{\mathcal{T}}(y|\tilde{\theta}_0)$$

- iii). If $K_{\epsilon}(\|s(y) - s(y_{obs})\|)$: move to *MCMC Step*.
 Else : return to i).

- *MCMC Step*.

- i). For $t \in [1, N]$:

- (a) Perturb the previous parameter value $\tilde{\theta}_{t-1}$ to get a new value θ^* .

$$\theta^* = K^*(\tilde{\theta}_{t-1})$$

- (b) Observe y from the theorised model \mathcal{T} using these perturbed parameters θ^* .

$$y \sim f_{\mathcal{T}}(y|\tilde{\theta}_0)$$

- (c) If $K_{\epsilon}(\|s(y) - s(y_{obs})\|)$: Set $\tilde{\theta}_t = \theta^*$.
 Else : Set $\tilde{\theta}_t = \tilde{\theta}_{t-1}$.

- Produce a posterior from the accepted parameter values $\{\tilde{\theta}_0, \dots, \tilde{\theta}_N\}$

3 Semi-Automatic ABC

Definition 3.1 - *Semi-Automatic ABC*[6]

In *Semi-Automatic ABC* summary statistics are learnt from simulation, but the user still has to make choices around what transformation $\mathbf{f}(\cdot)$ of simulated data \mathbf{y} .

An application of *Semi-Automatic ABC* should perform better in a general setting than traditional *ABC*.

Proposition 3.1 - *Semi-Automatic ABC - Algorithm*

- i). Perform a pilot run of ABC^[16] to determine a training-region of non-negligible posterior mass.
- ii). for $t \in [1, M]$:
 - (a) Simulate parameters θ_t from our prior $\pi_0(\theta)$, with the prior truncated to the training-region determined in (i).
 - (b) Simulate results $\mathbf{y}_t \sim f(\mathbf{y}|\theta_t)$ using these parameters.
- iii). Use simulated data and parameter values to estimate summary statistics.^[17]
- iv). Run ABC with these estimated-summary statistics.

Remark 3.1 - *Step iii)*

In step iii) we have simulated data $\mathcal{D} := \{(\theta_1, \mathbf{y}_1), \dots, (\theta_M, \mathbf{y}_M)\}$ where $\theta_t \in \mathbb{R}^m$, $\mathbf{y}_t \in \mathbb{R}_t^n \forall t \in [1, M]$. We want to learn a transformation $f(\mathbf{y}_t)$ of the simulated data \mathbf{y}_t st the parameters θ_t can be learnt from the transformation.

^[16]We need to define some arbitrary summary-statistics for this.

^[17]Potential methods inc. linear-regression, lasso analysis, cross-correlation analysis.

A *Linear Regression* approach is to learn the pick a function $\mathbf{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ ^[18] which maps the simulated data \mathbf{x}_t to the same dimension as the simulated parameters (ie to m dimensions) and then the parameters $\boldsymbol{\beta}_0, \boldsymbol{\theta}_1$ which give the least total-error across the whole data set \mathcal{D} ^[19] for the following

$$\begin{aligned}\boldsymbol{\theta}_t &= \boldsymbol{\beta}_0 + \boldsymbol{\beta}_1 \cdot \mathbf{f}(\mathbf{y}_t) + \varepsilon_t \\ \Leftrightarrow [\boldsymbol{\theta}_t]_i &= [\boldsymbol{\beta}_0]_i + [\boldsymbol{\beta}_1]_i \cdot [\mathbf{f}(\mathbf{y}_t)]_i + [\varepsilon_t]_i\end{aligned}$$

Our estimate for the model parameters $\boldsymbol{\theta}$, given some data \mathbf{y} , is thus the fitted value

$$\hat{\boldsymbol{\theta}} = \mathbb{E}[\boldsymbol{\theta}|\mathbf{y}] = \hat{\boldsymbol{\beta}}_0 + \hat{\boldsymbol{\beta}}_1 \mathbf{f}(\mathbf{y})$$

The constant terms $\hat{\boldsymbol{\beta}}_0$ can be ignored as ABC only uses the distance between summary statistics (not their absolute value). This means our m summary statistics are the different dimensions of $\hat{\boldsymbol{\beta}}_1 \mathbf{f}(\cdot)$

Remark 3.2 - Choosing transformation $\mathbf{f}(\cdot)$

In **Remark 1.2** the user has to define how to transform the simulated results (ie define $\mathbf{f}(\cdot)$) and this choice will affect the set of summary statistics generated. It is easy to run this stage multiple times, using different transformations on the same data \mathcal{D} and then using standard model comparison procedures^[20] to determine which of the generate summary statistics are sufficient.

^[18]This transformation can actually map to any dimension but we prefer for it to be a lower dimension than the simulated data \mathbf{y} .

^[19]Generally least-square-error.

^[20]e.g. BIC, sufficiency

References

- [1] Mark A. Beaumont. Approximate bayesian computation. *Annual Review of Statistics and Its Application*, 6(1):379–403, 2019.
- [2] Richard G. Everitt and Paulina A. Rowińska. Delayed acceptance abc-smc. *Journal of Computational and Graphical Statistics*, 0(0):1–12, 2020.
- [3] Patrick Laub. Approximate bayesian computation: Introduction insurance examples.
- [4] Tina Toni, David Welch, Natalja Strelkowa, Andreas Ipsen, and Michael P.H Stumpf. Approximate bayesian computation scheme for parameter inference and model selection in dynamical systems. *Journal of The Royal Society Interface*, 6(31):187–202, 2009.
- [5] Florian Hartig, Marshal says:, Florian Hartig says:, Matias Francisco Hernandez Arellano says:, and Nori Says:. A simple approximate bayesian computation mcmc (abc-mcmc) in r, Jan 2014.
- [6] Paul Fearnhead and Dennis Prangle. Constructing summary statistics for approximate bayesian computation: semi-automatic approximate bayesian computation. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 74(3):419–474, 2012.